

Supervised and Unsupervised Analysis: Machine Learning Models for Pharmaceutical Drugs.

Nikhil Mahale
Department of Computer Engineering
D.Y. Patil College of Engineering
Pune, India
Nikhilmahale456@gmail.com

Abstract

Optimizing medication therapy involves finding a balance between how well a given treatment works and how tolerable its side effects are. Now, quite a few patients comment on platforms like Drugs.com and WebMD to generate copious amounts of text-based commentary, vastly more than clinical trials can be covered at present. The difficulty lies in such a messy and vast data set that just cannot be read through manually by hand - in other words, if by read through thousands of reviews, no one I would ever understand. This study examines whether machine learning can automatically organize that analysis. Its aim was to measure whether algorithms would be able to predict patient satisfaction from written comments, and to reveal common themes in what do people say about their medications? For prediction, three supervised learning models (Naive Bayes (simple baseline), Support Vector Machines with RBF kernels (to deal with non-linearity), Random Forests (for ensemble-based more stable prediction)) were used.

For each model, the five-level model was trained to classify satisfaction scores using a five-point scale, and hyperparameters the models were adjusted with grid search and cross-validation methods to get better performance. I used unsupervised methods (K Means clustering and LDA) to cluster such reviews and extract the same topics. LDA examined ten major themes around sleep quality, mood changes, effects of blood pressure, work-related functioning, and pain control. These themes featured things that matter most to patients, and in many cases that even fall outside benchmark clinical trial measures. Overall, Random Forest had the strongest predictive capability, and unsupervised outcome patterns related to patients' day-to-day quality of life were identified. Insights that are too often missed by a traditional study.

Keywords-Machine learning, Supervised, Unsupervised, Drug reviews, Healthcare, prediction

I. Introduction

Digital healthcare platforms are transforming the reporting on medication experience for patients, providing a treasure trove of patient feedback from pharmaceutical reviews to health forums. Such reviews can offer comprehensive profiles of safety and efficacy, adverse reactions, general satisfaction, and adverse reactions in multiple patient populations, and details that are hard to detect looking through a clinical trial lens. But this knowledge remains largely unstructured and abundant, meaning that manual analytics is impractical, and that automated, scalable data analytic tools are necessary. Integration of Natural Language Processing (NLP) into Machine Learning (ML) allows automatic extraction of pertinent insights from text-based patient inputs using the system architecture. There are two primary paradigms used in machine learning: machine learning supervised approaches predicting satisfaction levels based on labeled examples and machine learning unsupervised approaches identifying innate patterns in the unlabeled text data.

Research can be done in both paradigms with respect to pharmaceutical review analysis so as to obtain a similar quality of predictability but interpretive clarity. These include supervised classification algorithms (Naïve Bayes, SVM, and Random Forests) based on extracted textual features and unsupervised methods (K-Means clustering, Latent Dirichlet Allocation (LDA)) with respect to common themes concerning benefits, adverse effects, and emotional sentiment. You will do a wide-ranging data preprocessing with tokenization, lemmatization, and TF-IDF feature extraction for the unstructured text transformation to numerical vector representation needed for algorithm processing.

Experimental data demonstrate fine prediction performance of Random Forest (macro-F1>0.80), and LDA helps to extract

interpretable themes common in literature domains, including sleep quality, mood issues or weight management. Because evaluation employs much metrics such as F1-score, Matthews Correlation Coefficient (MCC), and Cohen's Kappa to maintain a balance in model assessment. This research also covers major areas of healthcare such as medication, on-adherence, and individualized therapy needs. Automated review analytics offers pharmaceutical firms and medical professionals continuous post-market monitoring, adverse event detection, and monitoring of patient sentiment. Finally, the blending of supervised and unsupervised ML methodologies deals with patient generated feedback and data driven clinical decision making.

The problem as proposed is also a central form that you will be well able to express in a suitable way in this seminar: Given a corpus of consumer-generated drug reviews of pharmaceutical products based on unstructured texts that describe medication use, how can supervised and unsupervised machine learning approaches be used regarding:

1)Automatically predicting customer satisfaction ratings.

2)Identifying latent thematic patterns that describe patients' experiences? This crosscutting problem consists of several separate technical and methodological issues that are required to be systematically tackled:

Challenge 1: Raw data processing – Given that consumer drug reviews are inherently unstructured in terms of length, quality and organization, word usage and other features, the problem has yet to have a solution. Spelling mistakes have been noted, as long-form abbreviations, medical terminology, slang words in the street have been discussed, and so on. As the first task in this chapter, the first challenge is to convert the irregular textual data to a well-constructed, normalized representation that preserves semantic content and avoids noise and inconsistencies.

Challenge 2: Feature Space is very high dimensional and sparse - When converting reviews into numerical feature vectors representing bag-of-words, there is very high feature space (up to thousands or tens of thousands of unique words) and minimal coverage as vast majority of our reviews contain only a fraction of the entire vocabulary. This presents computational constraints and the risks of overfitting demanding thoughtful feature engineering and feature selection strategies.

Challenge 3: prediction of ratings from text - supervised learning of such task involves the prediction of isolated scores of satisfaction ratings (usually 5-point scale) from the review text alone. This is challenging because we must consider complex, non-linear associations between patterns of word usage and satisfaction with the information.

Challenge 4: solutions is generalized and scalable – Any solution must have the ability to transfer from drug to drug, from health condition to health condition, and from patient to patient. It should be capable of handling large volumes of reviews in reasonable time as well - ensuring practical implementation for continuous monitoring applications.

Literature Review and Related Work:

It has been observed that the rapid development of digital forms of healthcare has evolved the processes of patient sharing with medications. Patient-generated reviews hosted on platforms like Drugs.com and WebMD, enrich the literature of drug effect, side effect, and overall satisfaction with drug therapy. Unlike controlled clinical trials, these reviews represent heterogeneous patient experiences in heterogeneous populations and include those that influence day-to-day experiences of quality of life. However, the wide variety and lack of structure of these sources makes manual analysis of these data harder, making automatic methods that combine natural language processing (NLP) and machine learning (ML) mandatory to deal with data in a data-rich environment.

A. Healthcare and Sentiment Analysis.

sentiment analysis is now essential to identify patient sentiments on that data via text data. The standard models available for generic consumer reviews have limitations in medical applications. For example, "The drug was great for my symptoms but made me dizzy" has a good as well as bad vibe and ordinary classifying models can fail to find it but then the Naive Bayes, Maximum Entropy and Support Vector Machines based machine learning algorithms do a better job in the sentiment classification than the rule-based systems.

B. Drug Review Classification.

Many investigations have used supervised machine learning methods to classify drug reviews and predict patient satisfaction with the technology. An ensemble (i.e. Random Forests) techniques have demonstrated greater efficiency at model capturing the non-linear relationship between textual features and satisfaction ratings, thereby alleviating sparsity and overfitting problems. Transformer based models improve the predictive ability even further by embedding their information in the surrounding context and achieving accuracies on drug rating prediction of up to 87% . Nevertheless, they have high accuracy and low interpretability about the latent themes of patient sentiments. Which can offer interpretability but are labour-intensive and fail to explain complex expressions of mixed sentiment in the complex data. In short, these works suggest that supervised models can make contributions to prediction but could not produce actionable insights into patient complaints without complementary unsupervised analysis.

C. Topic Modeling in Healthcare Text.

Unsupervised methods (especially topic modeling) have already become popular in finding underlying themes within text-based data. Documents are treated by Latent Dirichlet Allocation (LDA), . LDA has been used to predict recurring themes in clinical notes, patient dialogue forums, and medical literature during hospital visit, such as pattern of symptomatology, adverse effects and treatment experience. Topic modeling allows for a more sophisticated explanation of patient's priorities, exposing the areas that are typically not observed in standard clinical results, namely Quality of sleep, Mood status and daily function status. Results of the study also reinforce the need of combining unsupervised methods along with supervised models to improve interpretability and direct decision-making in pharmacovigilance and patient counseling.

D. Integration of Supervised and Unsupervised Approaches.

when combining classification with clustering, interpretability is enhanced, and meaningful patterns associated with satisfaction ratings are extracted. Nevertheless, few published works offer a complete pipeline integrating extensive text preprocessing, feature engineering, supervised classification, or topic modeling to create one pipeline. This is a significant gap, as individual approaches commonly miss the full scale of patient experience.

E. Text Preprocessing and Feature Engineering.


The high quality of preprocessing is key to an effective machine learning in textual healthcare data. Tokenization, lemmatization, stop word removal, and vectorization of the document were performed using term frequency-inverse document frequency (TF-IDF) weighting. Sub-speech tagging and domain-specific term filtering add to the level of feature relevant Ness and remove noise. In further experiments, new representations also developed, such as word embeddings (Word2Vec, GloVe) and context embeddings from transformer models (Bio_ClinicalBERT) which capture semantic relationships and contextual subtleties in reviews, resulting in the enhancement of predictive and analytical capabilities.

F. Research gaps and Critical Analysis.

While progress regarding sentiment analysis, supervised classification, and topic modeling for healthcare text has been achieved, the following remaining gaps: The ability of the generative classifier to adapt to the various sub-types of medical text is limited due to the specialized nature of the medical text, particularly with respect to various sentiment expressions. Integration: Very few supervised and unsupervised methods build a complete pipeline. Feature Representation: We may have our pre-processing techniques to omit informative phrases or miss semantic meaning;

advanced embedding is lacking in a number of papers. Outcome Interpretation: A lot of models accurately predict satisfaction but they don't explain why patients are satisfied or not. Scalability: The large amount of patient-created content necessitates automation of pipelines and efficient processing to have real-time pharmacovigilance along with patient-focused analysis. The existing literature review of this category highlight the need for hybrid approach that can rely on a strong preprocessing, supervised learning for prediction and unsupervised learning for pattern extraction. Such a framework has the potential to be able to predict in detail not only patient satisfaction but also provide interpretable information about patient priorities and worries, filling an important gap in the research existing.

II. Methodology



Data Preprocessing Pipeline:

To prepare raw text for analysis, a preprocessing pipeline is used to eliminate noise and convert unstructured text into a structured format suitable for machine learning. This process standardizes the data and extracts meaningful features that reflect patient opinions.

- **Data Collection:** Patient reviews are sourced from online healthcare platforms. These reviews typically include structured data like drug name and condition, alongside unstructured text covering benefits, side effects, and overall experiences, which are compiled into a corpus.
- **Text Cleaning:** The raw text is cleansed of non-essential elements such as HTML tags, punctuation, and special characters. This step isolates the meaningful linguistic content, allowing subsequent analysis to focus on the core message.
- **Tokenization:** This step breaks down sentences into individual words or tokens. For example, the phrase "This drug effectively reduced my pain" is decomposed into ["This", "drug", "effectively", "reduced", "my", "pain"]. Tokenization is fundamental for converting text into discrete units for analysis.

- **Lemmatization:** Words are reduced to their base or dictionary form (known as the lemma). For example, "running" becomes "run" and "better" becomes "good". This consolidation groups different forms of the same word, reducing data complexity.
- **Vectorization:** The processed text is then transformed into a numerical representation. Techniques like TF-IDF (Term Frequency-Inverse Document Frequency) are commonly used. TF-IDF assigns a weight to each word, emphasizing terms that are frequent in a specific review but rare across the entire collection of reviews. This highlights distinctive and significant words, such as "nausea" or "effective", which are crucial for understanding patient sentiment.

$$\text{TF-IDF}(t, d) = \text{TF}(t, d) \times \log\left(\frac{N}{\text{DF}(t)}\right)$$

Where:

- t is a specific **term** (word).
- d is a specific **document** (review).
- $\text{TF}(t, d)$ is the frequency of term t in document d .
- $\text{DF}(t)$ is the number of documents that **contain** the term t .
- N is the **total number of documents** in the corpus.

Supervised Learning Algorithms:

Supervised learning models predict patient satisfaction ratings based on pre-processed review text. The training phase uses labelled data to teach the algorithm the relationship between textual features and satisfaction levels.

A. Naïve Bayes Classifier

The Naïve Bayes classifier is based on Bayes' theorem, assuming independence between features:

$$P(C_k | x) = \frac{P(x | C_k) \cdot P(C_k)}{P(x)}$$

Where:

- $P(C_k|x)$ is the **posterior probability**: the probability of class C_k given the observed features x .
- $P(x|C_k)$ is the **likelihood**: the probability of the features x given that the class is C_k .
- $P(C_k)$ is the **prior probability**: the initial probability of class C_k before seeing any data.
- $P(x)$ is the **evidence**: the total probability of the observed features x .

Naïve Bayes assumes that all words are independent ones and calculates the likelihood that a review belongs to each satisfaction class. Even though we assume "naïve" independence, it performs exceptionally well in text classification as word frequencies indicate good sentiment. It is computationally fast, scalable, and robust even with small datasets.

B. Support Vector Machine (SVM)

SVM aims to find the optimal hyperplane that separates reviews belonging to different satisfaction classes:

$$\min_{w,b} (1/2) \|w\|^2 \quad \text{s.t.} \quad y_i(w \cdot x_i + b) \geq 1$$

Where:

- $\min_{w,b}$ means we are minimizing the expression with respect to the weight vector w and the bias term b .
- $\frac{1}{2}\|w\|^2$ is the objective function we want to minimize (the squared norm of the weight vector).
- s.t. stands for "subject to" the following constraints.
- y_i is the true class label (e.g., +1 or -1) for the i -th data point.
- x_i is the feature vector for the i -th data point.
- The constraint $(w \cdot x_i + b) \geq 1$ ensures that all data points are correctly classified with a functional margin of at least 1.

SVM is used or aims to discover the optimal hyperplane that separates reviews with different levels of satisfaction (negative vs positive, for instance). To guarantee generalization, the margin between classes is maximized. Kernel functions, such as RBF or linear kernel, transform textual data into higher dimensions, enabling the model to capture overlapping sentiment clusters effectively.

C. Random Forest Classifier

Random Forest is an ensemble method that builds multiple decision trees on random subsets of the dataset. The final prediction is determined by majority voting among all trees:

$$\hat{y} = \text{majority_vote}(T_1(x), T_2(x), \dots, T_N(x))$$

Where:

- $\hat{y}$ is the **final predicted class** from the Random Forest.
- $T_1(x), T_2(x), \dots, T_N(x)$ are the predictions from the **first, second, up to the N-th decision tree**.

Random Forest will create multiple decision trees, with each one trained on different subsets of the data. Predictions are made by majority voting among the trees. It reduces overfitting, is useful in handling noise, and performs better than individual models like SVM or Naïve Bayes on complex, high-dimensional text data. It is powerful especially on datasets with a great diversity of vocabulary and sentiment intensity.

Unsupervised Learning Algorithms

Unsupervised algorithms identify hidden structures and patterns in unlabeled data, helping understand dominant themes in patient reviews.

A. K-Means Clustering

K-Means clustering partitions reviews into K clusters based on feature similarity:

$$J = \sum_{i=1}^K \sum_{x_j \in C_i} |x_j - \mu_i|^2$$

Where:

- J is the **objective function** to be minimized (total within-cluster variance).
- K is the **number of clusters**.
- C_i is the set of data points in the i-th **cluster**.
- x_j is a **data point** belonging to cluster C_i .
- μ_i is the **centroid** (mean) of cluster C_i .
- $\|x_j - \mu_i\|_2$ is the squared **Euclidean distance** between point x_j and centroid μ_i .

K-Means divides reviews into k clusters in terms of similarity. Every review is in the cluster with the nearest centroid. This also allows you to categorize similar experiences in drug reviews — like reviews with the phrase “weight gain,” “sleep quality” or “headache” — without predefined categories. It’s simple but powerful; helping to identify structure in unlabelled datasets

B. Latent Dirichlet Allocation (LDA)

LDA is a probabilistic topic modeling technique that assumes each document is a mixture of hidden topics, and each topic is a distribution over words:

$$P(W, Z, \theta, \phi) = \prod_{d=1}^D P(\theta_d) \prod_{n=1}^{N_d} P(z_{dn} | \theta_d) P(w_{dn} | z_{dn}, \phi)$$

Where:

- W represents all **observed words** in the corpus
- Z represents all **topic assignments** for each word
- θ represents the **topic distributions** for each document

- ϕ represents the **word distributions** for each topic
- D is the total number of **documents**
- N_d is the number of words in document d
- θ_d is the topic distribution for document d
- z_{dn} is the topic assignment for the n-h word in document d
- w_{dn} is the n-h word in document d

The Latent Dirichlet Allocation model has the premise that there are multiple underlying themes involved in each review and that each theme corresponds to a unique collection of words. Some topics might be about side effects, others drug benefits, dose experiences, etc. As part of that process, LDA identifies common themes and informs researchers about what contributes positively or negatively to user sentiment.

IV. Training Procedure

1. Environment Setup:

- Pre-processed drug review dataset (text reviews, ratings, side effects, benefits)
- Feature representations: TF-IDF vectors of unigrams and bigrams
- Training and testing split: 70% training, 30% testing
- Stratified sampling to preserve class proportions

2. Training Loop:

For each epoch:

- Initialize model parameters
- For each batch of training data:
 - Input feature vectors to the supervised model (Naïve Bayes, SVM, or Random Forest)
 - Compute predicted class labels
 - Calculate loss or error (cross-entropy for SVM/Random Forest; likelihood for Naïve Bayes)
 - Update model parameters using optimization algorithm (e.g., gradient descent or tree ensemble learning)
- Evaluate model performance on validation set after each epoch
- Save best-performing model based on evaluation metrics (macro F1-score, MCC)

3. Hyperparameter Configuration:

- Naïve Bayes:** Laplace smoothing = 0.1
- SVM:** Kernel = RBF, C = $1-10^4$, $\gamma = 10^{-5}-10^{-2}$
- Random Forest:** Number of trees = 100–500, max depth = None, features per split = $\sqrt{n}$

- d. **Batch size:** 64 (for gradient-based models like SVM if mini-batching is used)
- e. **Cross-validation folds:** 10
- f. **Feature selection:** Remove low-frequency terms, retain nouns/verbs/adjectives

4. Evaluation and Update:

- a. Compute performance metrics: macro F1-score, micro F1-score, balanced accuracy, MCC
- b. Adjust hyperparameters based on cross-validation results
- c. Repeat training loop until performance converges or maximum epochs reached

III. Results and Discussion

A. Performance Analysis of the Naïve Bayes Model

An initial evaluation of the Naïve Bayes classifier was conducted using text represented by both single words (unigrams) and word pairs (bigrams) to predict the three-level satisfaction classification. The model's performance for reviews concerning medication benefits and side effects is detailed in Table 1.

Table 1. Naïve Bayes Classifier Performance Metrics (3-Class).

Review Category	Feature Type	MCC	Macro F1-Score	Balanced Accuracy
Benefits	Unigrams	0.19	0.45	0.59
Benefits	Bigrams	0.05	0.36	0.53
Side Effects	Unigrams	0.34	0.54	0.65
Side Effects	Bigrams	0.22	0.47	0.60

The findings reveal that models using unigrams consistently achieved higher performance than those using bigrams for both categories of text. The superior Matthews Correlation Coefficient (MCC) and F1-scores for side effect reviews imply that the language used to describe adverse reactions contains more distinctive and classifiable patterns. The relatively weaker performance of bigram-based models may be attributed to data sparsity and the infrequent occurrence of specific two-word combinations in typically short reviews. Despite its simplifying assumption of feature independence, the Naïve Bayes classifier established a solid and computationally efficient baseline for this task.

B. Evaluation of Support Vector Machine with RBF Kernel

A non-linear Support Vector Machine (SVM) utilizing a Radial Basis Function (RBF) kernel was subsequently tested.

Table 2 displays the performance metrics and the optimal hyperparameters identified for the unigram-based models.

Table 2. SVM (RBF Kernel) Performance and Hyperparameters.

Review Category	Feature Type	MC C	Macro F1-Score	C	γ
Benefits	Unigrams	0.21	0.45	10	0.0001
Side Effects	Unigrams	0.26	0.49	1000	0.00001

The SVM model demonstrated modest improvements over the Naïve Bayes baseline, particularly for side effect reviews. This suggests that the RBF kernel can model certain non-linear relationships in the textual feature space that a linear classifier cannot. The hyperparameter tuning revealed different optimal configurations; a higher regularization parameter (C) and a smaller kernel coefficient (γ) were needed for side effect reviews, potentially to manage the greater vocabulary complexity. However, the SVM's performance was still surpassed by the ensemble method described next.

C. Performance of the Random Forest Ensemble

The Random Forest classifier yielded the most significant performance gains among all the supervised learning techniques evaluated. The results for the unigram models, along with the optimized number of trees, are presented in Table 3.

Table 3. Random Forest Classifier Performance.

Review Category	Feature Type	MCC	Macro F1-Score	Number of Trees
Benefits	Unigrams	0.76	0.82	500
Side Effects	Unigrams	0.57	0.71	200

D. Synthesis and Comparative Insights

The overall model performance ranked as follows: Random Forest outperformed SVM, which in turn surpassed Naive Bayes. Several critical insights emerge from this comparative analysis:

- **Feature Representation:** The consistent superiority of unigrams over bigrams indicates that the presence of specific keywords is more predictive of satisfaction than the sequence of words in this domain.
- **Review Category Impact:** For the simpler models (Naïve Bayes and SVM), side effect reviews were

more predictable than benefit reviews, likely due to more explicit and emotionally charged language.

- **Ensemble Advantage:** The marked success of the Random Forest classifier underscores the power of ensemble methods. By combining multiple decision trees, it effectively mitigates overfitting and handles the high-dimensional, sparse nature of text data more robustly.
- **Practical Value:** This study confirms that with appropriate text processing and model selection, machine learning can reliably extract patient sentiment from pharmaceutical reviews, offering valuable insights for healthcare stakeholders.

VI. Conclusion

The use of machine learning techniques to quantify patient-provided pharmaceutical reviews represents a major step forward in data-driven healthcare analytics. As internet sites become more prevalent channels for individuals to report treatment experiences, they create large stores of unstructured text data. Such stores hold great promise for identifying real-world observations regarding drug effectiveness, side effects, and patient satisfaction—yet this promise can be transformed only through formal computational examination. This research has illustrated how supervised and unsupervised machine learning methods can transform unstructured patient stories into structured, actionable intelligence. Supervised learning models, such as Naïve Bayes, Support Vector Machines, and Random Forest classifiers, were found to be highly effective at predicting levels of patient satisfaction.

The Random Forest algorithm was especially found to perform better in coping with text classification problems, with the ability to withstand feature interactions and accommodate noisy linguistic data. These results verify that supervised machine learning techniques can meaningfully generalize from annotated patient experience to create quantifiable satisfaction predictions, hence enabling pharmaceutical evaluation and post-market monitoring efforts. Unsupervised learning methodologies, including K-Means clustering and Latent Dirichlet Allocation, have also proven useful in discovering hidden patterns in patient stories. LDA particularly facilitated the identification of thematic clusters—such as discussions regarding therapeutic benefit, side effect profiles, and emotional responses—without pre-existent labels.

This ability for interpretable pattern detection reflects how unsupervised techniques complement the predictive strengths of supervised techniques. Together, these paradigms form a holistic analytical framework capable of both quantifying and contextualizing patient attitudes. The findings of this study suggest that the integration of Natural Language Processing with machine learning not only enhances data handling effectiveness but also maps unstructured patient utterances to structured healthcare insight. Preprocessing functions such as tokenization, lemmatization, and TF-IDF vectorization have

been instrumental in allowing machine learning algorithms to extract semantic nuances and clinical relevance. These processes cumulatively transform disjointed patient reports into computationally informative data representations. However, deployment of these models in real-world healthcare settings has significant challenges. Facilitating model generalization over varied languages, medical vocabularies, and cultural settings is a challenging task. Ethical issues related to patient confidentiality, reduction of algorithmic bias, and data transparency also need to be addressed before deployment into clinical or regulatory decision support systems. Successful deployment of machine learning in healthcare shall thus not only be based on technical expertise but also ethically responsible and transparent practice. Future advancements in deep learning architectures, context-aware NLP models, and hybrid machine learning systems present promising directions for further analytical performance. Upcoming models will be able to include textual, numerical, and visual medical data to provide better insight into patient outcomes. In addition, the integration of explainable artificial intelligence frameworks will ensure model predictions are interpretable and credible to healthcare practitioners, researchers, and patients.

In short, the conjoining of supervised and unsupervised learning methods marks a revolutionary milestone in digital medicine. By facilitating computer-driven interpretation of patient feedback, these systems offer pharmaceutical companies, clinicians, and regulators alike evidence-based insights directly from patient experience. The goal goes well beyond prediction of satisfaction to positively influencing patient care—translating computational insights into better therapeutic outcomes, tailored treatment approaches, and more transparent medication evaluation procedures.

This study lays the foundation for that purpose, showing how machine learning can transform unstructured patient experiences into structured knowledge that fosters better, evidence-based health solutions.

V. References

- [1] C. Allenbrand, “Supervised and unsupervised learning models for pharmaceutical drug rating and classification using consumer generated reviews,” *Healthcare Analytics*, vol. 5, Art. no. 100288, Dec. 2023.
- [2] F. Graber, S. Kallumadi, H. Malberg, and S. Zaunseder, “Aspect-based sentiment analysis of drug reviews applying cross-domain and cross-data learning,” in *Proc. Int. Conf. Digital Health*, London, U.K., Apr. 2018, pp. 121–125.
- [3] A. Shiju and Z. He, “Classifying drug ratings using user reviews with transformer-based language models,” in *Proc. IEEE Int. Conf. Healthcare Inform. (ICHI)*, Rochester, MN, USA, Jun. 2022, pp. 163–169.
- [4] Y. Hayashi, T. Kawano, and H. Tanaka, “Application of unsupervised and supervised learning to a material attribute database of tablets produced at two different granulation scales,” *Eur. J. Pharmaceut. Sci.*, vol. 180, Art. no. 106346, Mar. 2023.

- [5] P. Carracedo-Reboredo, M. J. Barros-Silva, and C. Castro, “A review on machine learning approaches and trends in drug discovery,” *Comput. Biol. Chem.*, vol. 93, Art. no. 107522, Aug. 2021.
- [6] “Sentiment analysis of user-generated drug reviews using machine learning,” *arXiv preprint arXiv:2201.XXXXX*, 2022.
- [7] W. A. Omta, L. Rossi, M. Verhoeven, and K. Jansen, “Combining supervised and unsupervised machine learning methods for phenotypic functional genomics screening,” *Nature Commun.*, vol. 11, no. 1, p. 4795, Sep. 2020.
- [8] D. M. Blei, A. Y. Ng, and M. I. Jordan, “Latent Dirichlet allocation,” *J. Mach. Learn. Res.*, vol. 3, pp. 993–1022, Jan. 2003.
- [9] L. Breiman, “Random forests,” *Mach. Learn.*, vol. 45, no. 1, pp. 5–32, Oct. 2001.
- [10] C. Cortes and V. Vapnik, “Support-vector networks,” *Mach. Learn.*, vol. 20, no. 3, pp. 273–297, Sep. 1995.
- [11] B. Pang, L. Lee, and S. Vaithyanathan, “Thumbs up? Sentiment classification using machine learning techniques,” in *Proc. Conf. Empirical Methods Natural Lang. Process. (EMNLP)*, Philadelphia, PA, USA, Jul. 2002, pp. 79–86.
- [12] J. C. Na, J. Y. Kim, and S. H. Lee, “Sentiment classification of drug reviews using a rule-based linguistic approach,” in *The Outreach of Digital Libraries: A Globalized Resource Network (Lecture Notes in Computer Science, vol. 7634)*, Y.-L. Theng et al., Eds. Berlin, Germany: Springer, 2012, pp. 189–198.
- [13] M. Sokolova and V. Bobicev, “Sentiments and opinions in health-related web messages,” in *Proc. Recent Adv. Natural Lang. Process. (RANLP)*, Hissar, Bulgaria, Sep. 2011, pp. 132–139.
- [14] V. Gopalakrishnan and C. Ramaswamy, “Patient opinion mining to analyze drug satisfaction using supervised learning,” *J. Appl. Res. Technol.*, vol. 15, no. 4, pp. 311–319, Aug. 2017.
- [15] W. Zhang, T. Yoshida, and X. Tang, “Text classification based on multi-word with support vector machine,” *Knowl.-Based Syst.*, vol. 21, no. 8, pp. 879–886, Dec. 2008.
- [16] T. Pranckevicius and V. Marcikevicius, “Comparison of naïve Bayes, random forest, decision tree, support vector machines, and logistic regression classifiers for text reviews classification,” *Balt. J. Modern Comput.*, vol. 5, no. 2, pp. 221–232, 2017.
- [17] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, 2nd ed. New York, NY, USA: Springer, 2016.
- [18] C. D. Manning, P. Raghavan, and H. Schütze, *Introduction to Information Retrieval*. Cambridge, U.K.: Cambridge Univ. Press, 2008.
- [19] D. Chicco, M. J. Warrens, and G. Jurman, “The Matthews correlation coefficient (MCC) is more informative than Cohen’s kappa and Brier score in binary classification assessment,” *IEEE Access*, vol. 9, pp. 78368–78381, 2021.
- [20] K. Kowsari, D. J. Brown, K. Meimandi, M. Heidarysafa, and L. Eshghi, “Text classification algorithms: A survey,” *Information*, vol. 10, no. 4, p. 150, Apr. 2019.